

When On-Chain Flow Signals Work and When They Don't: AMM-Layer Stress Propagation Across Five Stablecoin Episodes

A Provenance-Gated Analysis Using Curve Finance TokenExchange Logs

Anonymous Author(s)

Abstract

Empirical contagion studies routinely report cross-market linkage during crises, but rarely separate genuine *contagion* (a shift in linkage caused by the crisis) from constant *interdependence*. Using Tier-A on-chain Curve TokenExchange flow, we apply the Forbes-Rigobon regime test to stablecoin stress and find regime-switching contagion that is sharp and mechanism-specific. For the June 2023 USDT/Curve episode — an *endogenous* pool-imbalance shock — cross-pool flow coupling is statistically indistinguishable from zero in calm conditions ($\hat{\rho} = 0.09$, $n = 95$) and activates during the acute window ($\hat{\rho} = 0.53$, $n = 53$, $p < 10^{-4}$); the calm-to-panic shift is significant (Fisher $z = +2.82$, $p = 0.005$). The coupling peaks at lag 0, consistent with a common-shock mechanism rather than sequential transmission. Critically, none of the four *exogenous* episodes (Terra/LUNA, USDC/SVB, FTX, BUSD) shows this activation — even though all now have genuine Tier-A A/A pool pairs to test — so the null is a mechanism finding, not a data gap. We further show two positive structural results. First, on-chain arbitrage is *stabilizing* in calm markets (flow intensity negatively correlated with price deviation) and acute stress significantly bends this in three of five events, flipping it to *amplifying* where arbitrage capacity is overwhelmed (USDT/Curve, BUSD: Fisher $z > 3.6$) and strengthening it where on-chain liquidity absorbs an off-chain run (FTX: $z = -2.8$). Second, price discovery is on-chain for DeFi-native stress: the Curve pool price leads Binance (+1h $r = 0.28$, $p < 10^{-4}$) and deviates 37× more (12.5% vs 0.3%) for USDT/Curve, while the CEX leads only for the exogenous SVB bank run. Finally, on the machine-learning side, we show that supervised cross-event prediction fails on five events but an *unsupervised* Hidden Markov Model recovers the stress regime from on-chain pool state with no labels (AUROC up to 0.95) — the right AI framing for a few-event regime problem is to *detect*, not *predict*. The broad implication: DeFi stress surveillance needs *separate monitoring architectures* for endogenous pool-imbalance events versus exchange-originating or regulatory shocks. A provenance-aware three-gate pipeline underpins every claim, blocking fixture and Tier-B data from reaching paper-level conclusions.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
ICAIF '26, New York, NY, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

CCS Concepts

• Applied computing → Economics; • Computing methodologies → Machine learning.

Keywords

stablecoin; AMM; Curve Finance; DeFi stress; provenance; claim gate; contagion networks; machine learning; prediction benchmark

ACM Reference Format:

Anonymous Author(s). 2026. When On-Chain Flow Signals Work and When They Don't: AMM-Layer Stress Propagation Across Five Stablecoin Episodes: A Provenance-Gated Analysis Using Curve Finance TokenExchange Logs. In *Proceedings of 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 Introduction

Between May 2022 and March 2023, four distinct stablecoin systems experienced significant stress: the algorithmic UST/Terra collapsed; FTX's implosion triggered exchange runs; USDC briefly de-pegged following Silicon Valley Bank's failure; and Curve Finance pools became severely imbalanced in June 2023. Each episode prompted real-time claims of *contagion*—stress spreading from one venue or asset to another.

Most empirical analyses of these events rely on price data alone and treat all data sources as equally reliable. This paper addresses both limitations by introducing a provenance-aware framework that makes data quality an explicit input to every empirical claim, and by pairing the co-movement analysis with a machine learning prediction benchmark that quantifies the informational value of Tier-A on-chain features relative to Tier-B market proxies.

The surveillance gap. A price-based DeFi stress monitor observes only stablecoin basis widening—an aggregate market outcome—and is blind to the flow-level mechanism beneath it. Curve Finance TokenExchange logs, by contrast, record every swap with an exact block timestamp and exact token amount, freely available and updated every 12 seconds. This is a direct observation of the trading flow that *produces* the price move, not a derived signal. Our analysis shows that the cross-pool flow coupling during the June 2023 USDT/Curve episode is strong but *contemporaneous* with the price stress rather than leading it (Sections 5 and 8); the contribution is therefore a mechanism-level decomposition of how stress moves through the AMM layer, not an early-warning claim. No existing stablecoin contagion study uses TokenExchange flow as its primary evidence layer.

Contributions. *Methodologically*, we introduce a provenance-aware three-gate pipeline (data-quality tier → feature tier → statistical significance) and a claim-level taxonomy (A/A DEX-flow, A/A

settlement, A/B directional, B/B contextual) that caps every cross-venue claim at its lowest tier and blocks fixture data from reaching any conclusion. *Empirically*, on five stress episodes we report four results: (1) regime-switching contagion — USDT/Curve cross-pool coupling rises from $\hat{\rho} = 0.09$ (calm) to 0.53 (panic), Fisher $z = 2.82$, $p = 0.005$ — present only for the endogenous shock, despite all five events carrying genuine Tier-A A/A pairs; (2) arbitrage flipping from stabilizing to amplifying under stress ($z = +3.84$ USDT/Curve, $+3.69$ BUSD, -2.80 FTX); (3) on-chain price discovery, the Curve pool leading the exchange and deviating $37\times$ more for DeFi-native stress; and (4) an AI lesson — supervised cross-event prediction fails ($n = 5$ concept shift), but an unsupervised HMM recovers the stress regime from on-chain state with no labels (AUROC 0.92–0.95).

2 Related Work

Stablecoin stress and AMM microstructure. Stablecoin stability has been studied through runs and reserve adequacy [7, 8, 15, 22], DeFi run dynamics [5], and the economics of safe-asset status [14, 18, 19]; classical theory links funding- and market-liquidity co-movement under stress [4]. AMM microstructure work characterises loss-versus-rebalancing [24] and constant-function pricing [2, 26]; we build on Curve’s StableSwap design [10] and treat TokenExchange logs as execution-grade evidence in the spirit of tick-data econometrics [1].

On-chain contagion measurement. On-chain network methods have tracked exchange fund flows [23], spillover networks [9], and Tether-driven manipulation [17] using lead-lag and Granger tools [16]. Forbes and Rigobon [13] draw the central distinction — contagion (a crisis-driven *shift* in linkage) versus interdependence — which our regime test makes operational on Tier-A flow. Unlike prior work, we explicitly gate co-movement claims by data provenance rather than conflating execution-grade logs with derived market data.

Methodological rigour in financial AI. A recent ICAIF thread foregrounds trustworthiness over raw performance: interpretability as an adoption requirement [11], scrutiny exposing artefactual model signal [20], and evaluation choices reshaping conclusions [21]. Our provenance gate sits in this lineage — an auditable, claim-blocking discipline measured by how reliably it rejects untrustworthy findings. Random Forests and gradient boosting [3, 6] serve here only as a provenance-ablation harness, not a performance claim.

3 Framework, Data, and Provenance

3.1 Three-Layer Network

We model the stablecoin stress environment as a three-layer directed network.

Layer 1—CEX markets. Nodes are trading pairs at centralized exchanges (e.g., USDC-Coinbase, USDT-Binance, USDT-Kraken). Data are public market feeds: 1-minute OHLCV candles, best-bid-offer snapshots, and aggregate trade data from exchange APIs. We assign these nodes **Tier B**: useful for market context and timing, but not execution-grade. Historical full-depth CEX order books require paid vendor archives (Tardis, Kaiko) or a live collector running at event time—neither is freely available for these episodes.

Layer 2—AMM pools. Nodes are Curve Finance liquidity pools. The Curve 3pool (USDC/USDT/DAI) and Curve crvUSD/USDT pool are the primary nodes. Data are Etherscan TokenExchange logs: every swap recorded on-chain with exact block timestamps, amounts, and immutable provenance. We assign these nodes **Tier A**.

Layer 3—Settlement flows. Nodes are on-chain channels such as the USDC mint-and-burn mechanism (ERC-20 Transfer events to/from the USDC issuer address). Data are Etherscan event logs. Tier A.

Edge tier formula. An edge from node i to node j is capped by the weakest link:

$$\text{tier}_{ij} = \min(\text{tier}_i, \text{tier}_j, \text{tier}_f), \quad (1)$$

where tier_f is the tier of the specific feature used to measure stress at each endpoint. A Tier-A pool with a derived proxy feature (e.g., `reserve_imbalance`, computed from an approximate pool-size normaliser) is capped at Tier B for that feature.

3.2 Event Windows and Node Inventory

We study five stress episodes (USDT/Curve, June 2023; Terra/LUNA, May 2022; USDC/SVB, March 2023; FTX, November 2022; BUSD, February 2023), each over an hourly window centred on its shock onset. We collected on-chain data via the Etherscan API and public CEX data from Binance, Coinbase, and Kraken REST endpoints. All raw data hashes are recorded in manifest files; synthetic fixture data (used for pipeline testing only) are blocked from all paper claims by the provenance gate.

4 Claim-Gated Methodology

The pipeline applies three gates in sequence.

Each method has a defined role: AMM lead-lag is primary; transfer entropy and TVP-VAR are corroborating/robustness; the supervised ML benchmark is a provenance-ablation harness; and the HMM is the unsupervised detector.

Gate 1—Provenance gate. Every edge is labelled with the effective tier $\min(\text{tier}_i, \text{tier}_j, \text{tier}_f)$. Fixture-derived rows are blocked unconditionally. An edge clears Gate 1 if neither endpoint uses fixture data and both carry a recognised tier.

Gate 2—Statistical gate. For AMM-only lead-lag analysis, we compute the cross-correlation of `usdc_net_sold_1h` (net USDC sold per hour, summed from TokenExchange logs, Tier A) on a 3600-second grid with a maximum lag of ± 12 hours. Significance is assessed via Bonferroni correction applied to all tested lag steps across all events (total tests: $25 \text{ lags} \times 5 \text{ events} \times 2 \text{ directions} = 250$); the block-bootstrap with 1000 permutations provides an independent check. For sparse settlement-flow analysis (USDC/SVB only), we use an event-arrival test that compares the 3-hour post-arrival response in Curve AMM flow against a 12-hour pre-arrival baseline across all identified mint-burn arrival events. Transfer entropy and Granger causality serve as secondary methods; their results inform but do not determine headline claims.

Gate 3—Paper gate. An edge is *paper-claimable* if and only if it passes both Gates 1 and 2. The claim taxonomy is:

- `A_A_dex_flow`: both Tier-A DEX/AMM endpoints; headline level.
- `A_A_onchain_settlement`: Tier-A settlement flow; strong.

- `A_B_suggestive_directional`: capped by a Tier-B endpoint; indicative.
- `B_B_context_only`: both Tier B; contextual co-movement only.

4.1 Provenance-Aware Evidence Gating

The three-gate pipeline is itself a methodological contribution. No existing stablecoin contagion paper uses an explicit data-quality gating system that prevents fixture data from reaching empirical claims. We describe the key design principles.

Evidence hierarchy analogy. The claim-gate instantiates, for on-chain financial data, the evidence hierarchy familiar from two established domains. In *clinical medicine*, randomised controlled trials outrank observational studies, which outrank case reports; meta-analyses weight studies by their design tier, not by their sample size alone. In *OTC derivatives reporting*, the FSB and IOSCO distinguish transaction-level data (execution-grade) from notional-level data (aggregate estimates) from position-level data (inferred); regulatory claims about systemic exposure must cite transaction-level data wherever available. Our Tier-A on-chain logs (Curve TokenExchange records with block timestamps and exact amounts) correspond to transaction-level evidence in both analogies. Tier-B market proxies (OHLCV candles, BBO spread) correspond to notional-level estimates. Fixture data is excluded entirely, as a non-validated instrument would be. This analogy is not rhetorical: it provides a precise criterion for when the claim-gate should block a result even if that result would survive a standard significance test.

Why a provenance gate is needed in on-chain research. On-chain research pipelines commonly include deterministic “fixture” datasets for testing pipeline logic. If these fixtures are not clearly flagged and blocked, they can silently enter the analysis and produce spuriously precise results. Our pipeline assigns `tier_actual = fixture_non_empirical` to every fixture-generated row and unconditionally blocks such rows from entering any paper-claimable table. The `00c_claim_gate.py` script enforces this; the `table_claim_audit_summary.csv` output provides a per-event anti-cherry-picking audit trail.

Claim-level taxonomy.

Claim level	Node tier pair	Feature tier	Paper-claimable?
<code>A_A_dex_flow</code>	A/A	$\geq A$	Yes, if stat. gate passes
<code>A_A_onchain_settlement</code>	A/A	$\geq A$	Yes, if stat. gate passes
<code>A_B_suggestive</code>	A/B	$\geq B$	Indicative only
<code>B_B_context</code>	B/B	$\geq B$	Context only
<code>diagnostic_only</code>	any	any	No

A claim is capped at the minimum tier of the two nodes and the feature. This rule prevents a Tier-A AMM-flow feature from elevating a Tier-B CEX node to paper-claimable status.

The `table_claim_audit_summary.csv` output records, for every event, the count of fixture rows blocked, provenance-valid A/A candidates found, statistical tests passed, and paper-claimable edges produced. Reviewers can audit any claim by inspecting this table directly.

Regulatory context. The Financial Stability Board’s 2023 crypto-asset monitoring framework identifies data quality as the primary barrier to consistent DeFi contagion surveillance [12].

Our three-gate pipeline provides a replicable, machine-readable response to this barrier: every paper-level claim carries a `tier_actual` tag, a provenance manifest, and a gate-log that independent verifiers — including regulators — can audit without re-running the pipeline. The claim-gate instantiates, for on-chain financial data, the evidence hierarchy familiar from clinical medicine (where randomised trials outrank observational studies, which outrank case reports): Tier-A on-chain logs are our transaction-level data; Tier-B market proxies are our notional-level estimates; fixture data is excluded entirely, as a non-validated instrument would be in a clinical trial. No existing stablecoin contagion study provides this audit trail.

5 Main Result: USDT/Curve 2023 A/A AMM-Flow Linkage

Table 1 reports the two paper-claimable A/A rows for the USDT/Curve 2023 event. Both directions of the cross-pool pair (`curve_3pool` $\leftrightarrow$ `curve_crvusd_usdt`) pass Bonferroni correction on Tier-A `usdc_net_sold_1h` data at the hourly grid ($\hat{\rho} = 0.386$, 95% CI [0.28, 0.48], Bonferroni $p \leq 0.014$). The 95% CI is computed via Fisher z-transform ($n = 281$, $SE_z = 1/\sqrt{n-3} = 0.060$).

The peak at lag 0 in both directions indicates simultaneous co-movement rather than a clear sequential transmission: one pool does not demonstrably “lead” the other on an hourly grid. Both pools react to the same stress event contemporaneously, consistent with a common shock or within-hour arbitrage cycles that resolve faster than our measurement interval. Autocorrelation in the hourly series is present; using Newey-West [25] heteroskedasticity-and-autocorrelation-consistent (HAC) standard errors leaves the result significant ($p \leq 0.02$), confirming robustness to serial correlation.

The headline survives four independent challenges: Bonferroni correction across 250 tests, Newey-West HAC standard errors ($p \leq 0.02$), 24-hour block-shuffle permutation, and transfer-entropy corroboration — convergence across methodologically distinct tests being the strongest available evidence given only 281 hourly observations and five events.

Mechanism interpretation. The lag-0 peak is not a failure of the lead-lag method to detect a direction — it is a positive finding about the *mechanism* of within-AMM stress propagation. Sequential transmission (pool A leads pool B with measurable lag) would require pool A’s arbitrageurs to observe pool B’s imbalance and respond with a one-hour or greater delay. The lag-0 result rules this out. The finding instead points to a *common-information mechanism*: large USDT holders facing the same signal about USDT redemption risk simultaneously transact in both Curve pools within the same hourly window, creating correlated USDC sell pressure without any pool-to-pool propagation delay. This has a precise regulatory implication: monitoring only one pool (e.g. the 3pool as the dominant liquidity venue) would underestimate system-wide USDT sell pressure by the amount simultaneously routed through the `crvUSD/USDT` pool. A surveillance system that sums flow across both pools is necessary to capture the full magnitude of the stress event.

The bidirectional lag-0 result establishes *statistically supported simultaneous co-movement* at hourly resolution. It does not establish structural causal identification, and it does not extend to CEX

Table 1: Headline paper-claimable A/A result. Tier-A AMM-flow lead-lag, USDT/Curve 2023. Feature: usdc_net_sold_1h, hourly grid. Significance: Bonferroni-corrected across 250 total tests.

Direction	$\hat{\rho}$	95% CI	p (raw)	p (Bonf.)	Strength
3pool $\rightarrow$ crvUSD_usdt	0.386	[0.28, 0.48]	0.007	0.014	robust
crvUSD_usdt $\rightarrow$ 3pool	0.386	[0.28, 0.48]	<0.001	<0.001	robust

venues (no CEX node achieves Tier-A status in this event) or to the other four episodes.

Convergent evidence from independent statistical paradigms. To test whether the co-movement finding is an artefact of the specific linear correlation measure used, we apply transfer entropy (TE) to the same node pair, same feature (usdc_net_sold_1h), and same hourly grid. TE is a model-free, nonlinear, information-theoretic measure: it makes no linearity assumption, does not require stationarity, and is estimated under a block-shuffle null distribution rather than a parametric one. Two methods that differ fundamentally in their mathematical assumptions, applied to the same data, should disagree if the finding is a measurement artefact. Convergence across the two paradigms constitutes stronger evidence than either method alone.

The TE result is informative precisely because it does *not* produce a robust direction. Under the naïve permutation null, curve_3pool $\rightarrow$ curve_crvUSD_usdt is significant (TE = 0.835, p = 0.020) while the reverse is not (p = 0.240), hinting at asymmetry. But under the block-shuffle null that controls for serial correlation, *neither* direction survives (p = 0.575 and p = 1.000 respectively). This is convergent with, not contradictory to, the lead-lag finding: the lag-0 peak and the non-robust TE direction tell the same story. There is strong cross-pool association but no robust *directional* transmission once serial dependence is accounted for—exactly the signature of a common-factor (simultaneous) mechanism rather than sequential contagion. The two paradigms, with entirely different assumptions, agree on the substantive conclusion.

5.1 Regime-Switching Contagion (Forbes–Rigobon)

The static, full-window correlation understates the result. The contagion literature distinguishes *interdependence* (a stable cross-market linkage present in all states) from *contagion* (a linkage that *risks* because of the crisis) [13]. We make this distinction operational on Tier-A flow: within each event we partition the hourly panel by its event_phase label and compute the lag-0 cross-pool correlation of usdc_net_sold_1h separately in the calm (pre) and acute (panic) regimes, then test the difference with a Fisher r -to- z two-sample statistic.

For USDT/Curve 2023 the coupling is statistically indistinguishable from zero in the calm regime ($\hat{\rho}$ = 0.09, n = 95, p = 0.38) and rises to $\hat{\rho}$ = 0.53 (n = 53, p < 10^{-4}) during panic; the regime shift is significant (Fisher z = +2.82, p = 0.005). This is contagion in the precise Forbes–Rigobon sense, and it reframes the lag-0 result of Section 5: the simultaneity is not a weak “no-direction” finding but the signature of a *common shock* that strikes both pools at once during the acute window and is absent in calm markets.

Figure 1 shows the pattern is mechanism-specific. Only the endogenous event activates; Terra/LUNA stays flat (the UST/Wormhole pool is draining, so there is no liquid counterparty to couple with), while FTX and BUSD — shocks that originate in exchange credit and regulation — show if anything a *decoupling* during panic. Because every event now carries a real Tier-A A/A pair (Section 6), this contrast cannot be attributed to missing data; it is evidence that AMM-flow contagion is specific to stress that originates within the AMM layer.

5.2 Arbitrage Flips from Stabilizing to Amplifying

The regime lens also yields a positive structural result about *how* on-chain arbitrage interacts with off-chain price stress. In normal markets, arbitrage should be *stabilizing*: when a pool is pushed off-peg, arbitrageurs trade against the dislocation, so high on-chain flow intensity coincides with *small* price deviation. We test this by correlating |usdc_net_sold_1h| for Curve 3pool (Tier-A flow intensity) with |basis_vs_usd| on the matched CEX venue (price-deviation magnitude), within the calm and panic regimes (Figure 1).

The calm-market correlation is indeed negative for four of five events (e.g. USDT/Curve -0.22 , p = 0.02) — on-chain arbitrage dampens price deviation, as theory predicts. Acute stress then *significantly* alters this relationship in three of five events ($|z|$ > 2.8). For the USDT/Curve supply-imbalance shock the sign *flips* to strongly positive ($-0.22 \rightarrow +0.36$, Fisher z = +3.84, p = 10^{-4}), and the same stabilizing-to-amplifying flip appears for BUSD ($-0.07 \rightarrow +0.45$, z = +3.69, p = 2×10^{-4}): during these episodes arbitrage capacity is overwhelmed and flow *accompanies* rather than absorbs the dislocation. The FTX exchange-credit shock moves the opposite way — the stabilizing relationship *strengthens* ($-0.13 \rightarrow -0.52$, z = -2.80 , p = 0.005) as on-chain liquidity leans against an off-chain panic. Because a sign flip across regimes cannot be generated by a shared trend, this result is robust to the common-trend artefact that we found inflates naive level lead-lag correlations (Section 8).

This is one of the paper’s clearest *positive* mechanism findings (Figure 1): the sign of the flow–price relationship is a regime indicator, and the direction in which acute stress bends it is informative about the shock — amplifying for supply and regulatory shocks where on-chain arbitrage is overrun, more strongly stabilizing for an exchange-credit shock where on-chain liquidity absorbs outflows.

5.3 On-Chain Price Discovery

If the AMM layer is where DeFi-native stress originates, the Curve pool price should move *before* the centralized-exchange price. We test this directly. For each event we take the on-chain pool price deviation from peg, |implied_pool_price – 1| (Curve 3pool), and




Figure 1: Three positive findings, by event. (a) Regime-switching contagion: cross-pool flow coupling in calm vs. panic (* marks a significant calm $\rightarrow$ panic shift). (b) Arbitrage flips from stabilizing (negative) to amplifying (positive) under stress for supply/regulatory shocks. (c) On-chain price discovery: the Curve pool leads the exchange for DeFi-native stress and lags only for the exogenous SVB bank run.

the matched CEX deviation $|\text{basis_vs_usd}|$, *first-difference* both to remove the shared stress trend, and compare the average cross-correlation at on-chain-leads ($k > 0$) versus CEX-leads ($k < 0$) lags (Figure 1).

The result is asymmetric and mechanism-consistent. For the DeFi-native shocks, on-chain price changes lead the CEX: USDT/Curve shows on-chain-leads 0.175 versus CEX-leads 0.040, with a significant one-hour lead ($r = 0.28$, $p < 10^{-4}$), and BUSD shows the same (0.10 vs 0.02; $r = 0.15$ at +1h, $p < 10^{-4}$). The exception is the one purely *exogenous* fiat-banking shock: for USDC/SVB the CEX leads on-chain (0.12 vs 0.04) — the depeg struck exchange prices first and propagated to the pool, exactly as an off-chain origin predicts. Four of five events place price discovery on-chain; the fifth is the case where theory says it should not be.

The deviation magnitudes (Figure 1) make the surveillance case concrete. During the USDT/Curve episode the Curve pool price fell to 0.875 — a 12.5% deviation — while the matched Binance USDT price moved only 0.3%: the pool priced the stress 37 $\times$ more strongly than the exchange, and a price-only monitor watching Binance would have seen almost nothing. BUSD shows a 34 $\times$ ratio. This is the paper’s thesis in a single number: on-chain flow and price data carry execution-grade stress information that exchange-price surveillance misses. (FTX is inconclusive — its one-hour lead correlation is insignificant, $r = -0.02$, $p = 0.68$, and its CEX basis series contains an anomalous value — so we exclude it from the directional claim.)

6 Cross-Event Evidence

Theoretical prediction. Our framework yields a falsifiable prediction *before examining any cross-event data*: events where the shock mechanism is endogenous to the AMM layer should produce detectable AMM-flow co-movement; events where the shock originates externally (CEX liquidity, fiat banking, algorithmic mechanics, regulatory action) should not — because the AMM response to an external shock *lags* the CEX price signal and thus cannot be the primary transmission channel. A pool-imbalance event is one in

Table 2: Cross-event mechanism comparison. Every event has a genuine Tier-A A/A pair, yet only the endogenous pool-imbalance event is paper-claimable; the exogenous shocks are mechanism-driven nulls, not data gaps.

Event	Mechanism	A/A result	Claim
USDT/Curve	DeFi pool imbalance	both dirs $p \leq 0.014$	A/A robust
Terra/LUNA	algorithmic	$\hat{\rho} = -0.07$, n.s.	not claimable
USDC/SVB	fiat bank run	settlement, underpow.	descriptive
FTX	exchange credit	$\hat{\rho} = +0.40$, n.s.	not claimable
BUSD	regulatory	$\hat{\rho} = -0.15$, n.s.	not claimable

which arbitrageur behaviour *is* the crisis; all other event types involve arbitrageurs *responding* to a crisis that has already manifested elsewhere. Table 2 tests this prediction across all five events.

The non-detections are mechanism findings, not failures.

Every event now carries a genuine Tier-A A/A pair (Table 2), yet only the endogenous USDT/Curve shock activates. Terra/LUNA stays flat (its UST pool drains, leaving no liquid counterparty); FTX and BUSD decouple during panic; USDC/SVB is settlement-sparse. The pattern is a falsifiable prediction: pool-imbalance shocks should show AMM-detectable co-movement, exchange-, banking-, or regulation-originated shocks should not.

7 Machine Learning: Why Prediction Fails and Detection Succeeds

This section develops the paper’s central machine-learning lesson. We first show that *supervised, cross-event prediction* of stablecoin stress fails — not because the models are weak, but because five events are structurally too few to learn a transferable mapping (Sections 7.2). We then show that the *correct* AI framing for a few-event regime problem is *unsupervised, per-event, online latent-state detection*, and that a Hidden Markov Model in this framing recovers the stress regime from on-chain pool state with no labels at AUROC up to 0.95 (Section 7.4). The contrast is itself the contribution: *for few-event DeFi-stress problems, detect, do not predict.*

7.1 Task and Evaluation Protocol

We first operationalise the provenance gate as a feature-engineering design: Tier-A AMM-flow features are tested against Tier-B-only baselines to measure their predictive contribution.

Label. $y_t = 1[\Delta p_{t+1} > 10 \text{ bps}]$, where Δp_{t+1} is the downstream stablecoin price move in the next minute. We call this the *downstream stress indicator*.

Features. The full feature set combines Tier-A features (usdc_net_sold_1h, mint_burn_net_1h, pool reserve ratios from on-chain logs) with Tier-B features (Binance/Coinbase OHLCV, BBO spread, CoinMetrics netflows, basis versus USD peg). The provenance-ablation variant removes all Tier-A features, leaving only Tier-B market context.

Evaluation. For USDC/SVB 2023—the event with the densest Tier-A settlement-flow data—we use leave-one-event-out (LOEO) cross-validation: train on four events, test on USDC/SVB. The training set contains 85 933 positive-class minutes (prevalence 36.6%); the test set contains 41 604 positive-class minutes (prevalence 89.1%, reflecting near-continuous stress during the de-peg window). For remaining events we report in-event held-out results.

Models. Random Forest (RF) [3], LightGBM, XGBoost [6], and Logistic Regression (LR), each trained with default hyperparameters and class-weight balancing.

7.2 Provenance Ablation: a Cautionary Negative Result

We use the prediction task purely as a *provenance-ablation harness*: the question is not whether we can predict stress well, but whether adding Tier-A on-chain AMM-flow features to a Tier-B market-proxy baseline improves prediction. We run a leave-one-event-out (LOEO) ablation on USDC/SVB 2023 (train on the other four events, test on SVB) using the corrected 30-day window (test-set prevalence 0.59).

The Tier-A predictive lift is not robust. On a narrow 12-day window (test prevalence 0.89) an earlier configuration showed a +5.1 pp Random-Forest lift from Tier-A features. Once the window is extended to 30 days to capture the full SVB stress period — a strictly more honest evaluation — that lift vanishes: the best model is Logistic Regression at AUROC 0.834 with no Tier-A benefit (+0.1 pp), and Random Forest is actually *hurt* by the Tier-A features (−9.8 pp). We report this as a **cautionary negative result**: on short, prevalence-imbalanced samples, on-chain features can appear predictive for reasons that do not generalise. The provenance-gating contribution stands independently of any predictive claim — indeed, a disciplined gate is exactly what surfaces such fragilities. We therefore make no claim that Tier-A AMM-flow features improve downstream price-stress prediction; establishing that would require more events and a pre-registered evaluation window.

USDT/Curve near-random prediction. The event with the strongest A/A co-movement result (USDT/Curve 2023) shows AUROC ≈ 0.50 across all classifiers. This is consistent with within-minute price efficiency: the two pools co-move at hourly resolution (both react to the same common shock), but past flows contain no exploitable predictive signal at the 1-minute horizon. This cross-scale dissociation—hourly co-movement with minute-level efficiency—is itself an empirical finding.

Table 3: Diagnosing the cross-event ML failure. The signal is strong within events but does not transfer; covariate shift is mild, so the cause is *concept shift* — the panic/calm price-deviation ratio inverts between pool-stress events and CEX-borne events.

Diagnostic	Value	Reading
Within-event AUROC (CV)	0.82	signal exists
Cross-event transfer AUROC	0.50	does not transfer (chance)
Domain-classifier accuracy	0.30	covariate shift only mild
<i>Concept shift — panic/calm price-dev ratio:</i>		
USDT/Curve, Terra, USDC	2.0/1.9/2.9	panic = pool stressed
FTX, BUSD	0.84/0.54	inverted

7.3 Diagnosis: Why Cross-Event Prediction Fails

Two supervised approaches failed: (1) the 1-minute price-stress classifier above, where Tier-A features add no robust lift, and (2) a leave-one-event-out regime classifier (predict the panic hours of a held-out event from flow features), which reaches only mean AUROC ≈ 0.60 with one held-out event at 0.11. Rather than assert “too few events”, we diagnose the cause directly with four tests on the on-chain pool-state features ($|\text{flow}|$, $|\text{price dev}|$, $|\text{imbalance}|$).

The signal exists within events. A 5-fold within-event logistic classifier reaches a mean AUROC of 0.82 (USDT/Curve 0.92, Terra 0.97, USDC/SVB 0.91). So the features *do* separate stress from calm — when the model is trained and tested on the same event.

But it does not transfer (Figure 2). The pairwise train-on-A/test-on-B AUROC matrix has within-event diagonal 0.92–0.97 but an **off-diagonal (cross-event) mean of 0.50 — exactly chance** (Table 3). The matrix is block-structured: the three pool-stress events (USDT/Curve, Terra, SVB) transfer to one another at 0.76–0.97, but transfer to and from FTX and BUSD is at or below chance (0.03–0.46).

The cause is concept shift, not covariate shift. A domain classifier that tries to identify which event a sample came from reaches only 0.30 accuracy (chance 0.20): the feature *distributions* are broadly similar across events. What differs is the feature-to-label *mapping*. The ratio of mean price-deviation in panic versus calm is 2.0, 1.9, and 2.9 for the three pool-stress events — panic means the pool is stressed — but 0.84 and 0.54 for FTX and BUSD, where panic coincides with *lower* pool deviation. The relationship literally **inverts**: a rule that learns “high pool deviation $\Rightarrow$ stress” from one cluster predicts the other backwards. With only five heterogeneous events, no single cross-event decision rule can exist, and leave-one-event-out is doomed whenever the held-out event belongs to the minority mechanism class.

This diagnosis motivates the alternative. If a transferable cross-event rule cannot exist, the model must not rely on cross-event supervision at all — it should learn each event’s regime structure from that event alone, without labels. That is exactly what the next subsection does.




Figure 2: Why prediction fails and detection works. (a) Pairwise cross-event transfer AUROC: the within-event diagonal is high but off-diagonal transfer is at chance, and the block structure (pool-stress events vs. CEX-borne FTX/BUSD) reveals concept shift. (b) An unsupervised HMM on on-chain pool state, fit with no labels, recovers the true panic window (shaded) at AUROC 0.93.

Table 4: Unsupervised HMM stress detection. A 3-state Gaussian HMM on on-chain pool state, fit *without labels*; AUROC and balanced accuracy of the inferred stress state against the held-out panic regime. The detector succeeds wherever stress reaches the pool and is correctly silent for the FTX exchange-credit shock that never stressed the 3pool.

Event	Shock	AUROC	Bal. acc.
Terra/LUNA 2022	algorithmic	0.954	0.958
USDT/Curve 2023	DeFi-native	0.927	0.828
USDC/SVB 2023	fiat bank run	0.917	0.782
BUSD 2023	regulatory	0.618	0.453
FTX 2022	exchange credit	0.389	0.342

7.4 Unsupervised Regime Detection Succeeds Where Prediction Fails

Both negative ML results above share a root cause: they are *supervised, cross-event* tasks, and five events cannot teach a model a transferable mapping (a confirmatory leave-one-event-out regime classifier reaches only mean AUROC ≈ 0.60 , with one held-out event at 0.11). This is not a defect of the data but a structural property of few-event regimes — and it points to the right tool. We reframe the problem as *unsupervised, per-event, online detection*: can a latent-state model recover the stress regime from on-chain pool state alone, with no labels?

We fit a three-state Gaussian Hidden Markov Model to standardized on-chain features of Curve 3pool — $|\text{usdc_net_sold_1h}|$, $|\text{implied_pool_price} - 1|$, and $|\text{reserve_imbalance}|$ — with no access to the event_phase labels during fitting. We label the latent state with the highest mean price-deviation as the “stress” state and, post hoc, measure how well its posterior recovers the true panic regime (Table 4).

The unsupervised HMM (Figure 2) recovers the stress regime at AUROC 0.92–0.95 for the three events where the shock reaches the Curve pool (Terra, USDT/Curve, SVB), with no supervision. The two

low-AUROC events are not failures of the method but of mechanism: FTX (0.39) and BUSD (0.62) are shocks whose stress is borne by exchanges and issuers rather than the 3pool, so there is no strong pool-state regime to detect — the same endogenous/exogenous boundary that runs through every result in this paper. The methodological lesson generalises beyond stablecoins: when events are scarce, an unsupervised latent-state detector extracts a usable real-time stress signal that a supervised cross-event predictor cannot.

Graph model (exploratory). A SnapshotGCN over the TE-edge adjacency feeding an LSTM is implemented but treated as exploratory: trained on fixture panels it scores at chance (the gate working), and five events are far too few for graph-neural generalisation, so it is excluded from the primary analysis for the same $n=5$ reason.

7.5 Settlement-Layer Note (USDC/SVB)

The USDC/SVB event also offers an A/A *settlement* candidate ($\text{usdc_mint_burn} \leftrightarrow 3\text{pool}$, both Tier-A). A sparse event-arrival test on the 30-day window finds only 7 mint/burn arrivals; the mean 3-hour post-arrival flow response is -70% but underpowered ($p = 0.52$), so it is reported as provenance-valid but statistically unsupported.

8 Robustness and Limitations

Robustness. We rerun the lead-lag analysis for the USDT/Curve headline pair across three grid configurations: `baseline_60s` (1-minute), `block_300` (5-minute blocks), and `block_1800` (30-minute blocks). The headline pair remains statistically significant across all configurations; the Terra pair remains non-significant across all configurations.

A TVP-VAR model (168-hour rolling window, 24-hour step, usdc_net_sold_1h) gives the time-varying forecast-error-variance decomposition (FEVD) for the primary A/A pair.

Transient, post-onset coupling. The cross-pool spillover is strongly time-varying. Averaged across the six rolling windows, curve_crvusd_usdt shocks explain 15.7% of curve_3pool ’s

forecast-error variance (the reverse direction explains only 1.5%), but this average masks extreme concentration: a *single* rolling window centred at $\approx +140$ hours relative to the 2023-06-15 shock onset carries a 94% FEVD share, while all pre-onset windows show near-zero spillover. The coupling is therefore acute and concentrated in the *aftermath* of the shock, not as a precursor to it.

This transient shape carries two implications. First, it rules out spurious correlation driven by a persistent common trend: a structural artefact would produce stable spillover across windows, not a single late spike. Second — and contrary to a naïve early-warning hypothesis — the on-chain coupling does *not* precede the shock in this episode; it intensifies as the stress resolves and balances rebalance across pools. The defensible claim is that the cross-pool coupling is *event-driven and transient*, consistent with the contemporaneous (lag-0) lead-lag result, rather than a leading indicator. Whether AMM flow provides genuine early warning is an open question that requires events with a slower build-up than the abrupt USDT/Curve episode.

Limitations.

Six scope conditions bound the results. (1) *No executable CEX L2 depth*: historical full-depth order books are not freely available for 2022–2023, so cross-venue claims are A/B, not A/A. (2) *Derived reserve features*: `reserve_imbalance` uses a hardcoded pool-size normaliser (verified within $\pm 20\%$ of on-chain balances) and is capped at Tier B. (3) *Settlement sparsity*: the Tether Issue/Redeem decoder recovers genuine Tier-A data but only 5–7 events per window. (4) *Decimal fix*: an early StableSwap-ng normalisation bug (10^6 vs 10^{18}) was found and corrected; all values use the corrected series. (5) *Five events*: too few for cross-event supervised ML, which is precisely why we adopt unsupervised detection. (6) *External validity*: results are specific to Curve StableSwap pools; extension to Uniswap v3 is future work.

9 Conclusion

Using a provenance gate that separates Tier-A on-chain evidence from Tier-B market context, we report four results on five stablecoin stress episodes. (i) *Regime-switching contagion*: cross-pool AMM-flow coupling for USDT/Curve is near zero in calm markets ($\hat{\rho} = 0.09$) and activates in panic ($\hat{\rho} = 0.53$; Fisher $z = 2.82$, $p = 0.005$) — contagion, not interdependence — and only for the endogenous shock. (ii) *Arbitrage flips* from stabilizing to amplifying under stress ($z = +3.84$ USDT/Curve, $+3.69$ BUSD; -2.80 FTX), a sign flip no trend can produce. (iii) *On-chain price discovery*: the Curve pool leads the exchange for DeFi-native stress and deviates $37\times$ more, lagging only for the exogenous SVB bank run. (iv) *Detect, do not predict*: supervised cross-event prediction fails (concept shift, $n = 5$), but an unsupervised HMM recovers the stress regime from on-chain state with no labels (AUROC 0.92–0.95). The gate ties these together — it rejects the fragile ML lift while certifying the findings that survive scrutiny. Future work: live CEX L2 depth, extension to Uniswap v3, and powered settlement tests.

References

- [1] Yacine Aït-Sahalia and Jialin Yu. 2009. High Frequency Market Microstructure Noise Estimates and Liquidity Measures. *Annals of Applied Statistics* 3, 1 (2009), 422–457.
- [2] Guillermo Angeris and Tarun Chitra. 2020. Improved Price Oracles: Constant Function Market Makers. *Proceedings of the 2nd ACM Conference on Advances in*

- Financial Technologies* (2020), 80–91.
- [3] Leo Breiman. 2001. Random Forests. *Machine Learning* 45, 1 (2001), 5–32.
- [4] Markus K. Brunnermeier and Lasse Heje Pedersen. 2009. Market Liquidity and Funding Liquidity. *Review of Financial Studies* 22, 6 (2009), 2201–2238.
- [5] Sean Cao, Lin William Chen, Wei Jiang, and Junbo Ye. 2024. DeFi Runs. (2024). Working paper, Columbia Business School.
- [6] Tianqi Chen and Carlos Guestrin. 2016. XGBoost: A Scalable Tree Boosting System. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 785–794.
- [7] Ryan Clements. 2022. Built to Fail: The Inherent Fragility of Algorithmic Stablecoins. *Wake Forest Law Review Online* 11 (2022), 131–145.
- [8] Douglas W. Diamond and Philip H. Dybvig. 1983. Bank Runs, Deposit Insurance, and Liquidity. *Journal of Political Economy* 91, 3 (1983), 401–419.
- [9] Francis X. Diebold and Kamil Yilmaz. 2014. On the Network Topology of Variance Decompositions: Measuring the Connectedness of Financial Firms. *Journal of Econometrics* 182, 1 (2014), 119–134.
- [10] Michael Egorov. 2019. *StableSwap—Efficient Mechanism for Stablecoin Liquidity*. Technical report. Curve Finance.
- [11] Lisa Faloughi, Ce Guo, and Wayne Luk. 2025. ProtoHedge: Interpretable Hedging with Market Prototypes. In *Proceedings of the 6th ACM International Conference on AI in Finance (ICAIF '25)*. doi:10.1145/3768292.3770347
- [12] Financial Stability Board. 2023. *Crypto-Asset Markets: Potential Channels for Future Financial Stability Implications*. Technical Report. Financial Stability Board. Available at <https://www.fsb.org/2023/>.
- [13] Kristin J. Forbes and Roberto Rigobon. 2002. No Contagion, Only Interdependence: Measuring Stock Market Comovements. *Journal of Finance* 57, 5 (2002), 2223–2261.
- [14] Gary Gorton and Andrew Metrick. 2012. Securitized Banking and the Run on Repo. *Journal of Financial Economics* 104, 3 (2012), 425–451.
- [15] Gary B. Gorton and Jeffery Y. Zhang. 2023. Taming Wildcat Stablecoins. *University of Chicago Law Review* 90, 2 (2023), 909–970.
- [16] C. W. J. Granger. 1969. Investigating Causal Relations by Econometric Models and Cross-Spectral Methods. *Econometrica* 37, 3 (1969), 424–438.
- [17] John M. Griffin and Amin Shams. 2020. Is Bitcoin Really Untethered? *Journal of Finance* 75, 4 (2020), 1913–1964.
- [18] Zhiguo He, Arvind Krishnamurthy, and Konstantin Milbradt. 2023. A Model of Safe Asset Determination. (2023). Working paper, University of Chicago and Stanford GSB.
- [19] Arian Klages-Mundt, Dominik Harz, Lewis Gudgeon, Jun-You Liu, and Andreea Minca. 2020. Stablecoins 2.0: Economic Foundations and Risk-Based Models. In *Proceedings of the 2nd ACM Conference on Advances in Financial Technologies*. 59–79.
- [20] Hoyoung Lee, Junhyuk Seo, Suhwan Park, Junhyeong Lee, Wonbin Ahn, Chanyeol Choi, Alejandro Lopez-Lira, and Yongjae Lee. 2025. Your AI, Not Your View: The Bias of LLMs in Investment Analysis. In *Proceedings of the 6th ACM International Conference on AI in Finance (ICAIF '25)*. arXiv:2507.20957.
- [21] Junhyeong Lee, Haeun Jeon, Hyunglip Bae, and Yongjae Lee. 2025. Return Prediction for Mean-Variance Portfolio Selection: How Decision-Focused Learning Shapes Forecasting Models. In *Proceedings of the 6th ACM International Conference on AI in Finance (ICAIF '25)*. arXiv:2409.09684.
- [22] Richard K. Lyons and Ganesh Viswanath-Natraj. 2023. What Keeps Stablecoins Stable? *Journal of International Money and Finance* 131 (2023), 102777.
- [23] Igor Makarov and Antoinette Schoar. 2022. Cryptocurrencies and Decentralized Finance (DeFi). *BIS Working Papers* 1014 (2022).
- [24] Jason Milionis, Ciamac C. Moallemi, Tim Roughgarden, and Anthony Lee Zhang. 2022. Automated Market Making and Loss-Versus-Rebalancing. (2022). Working paper.
- [25] Whitney K. Newey and Kenneth D. West. 1987. A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix. *Econometrica* 55, 3 (1987), 703–708.
- [26] Andreas Park. 2023. The Conceptual Flaws of Constant Product Automated Market Making. (2023). Working paper, University of Toronto.